

[Chapter 1 Introduction](#)

[Chapter 2 Operating-System Structures](#)

[Chapter 3 Process](#)

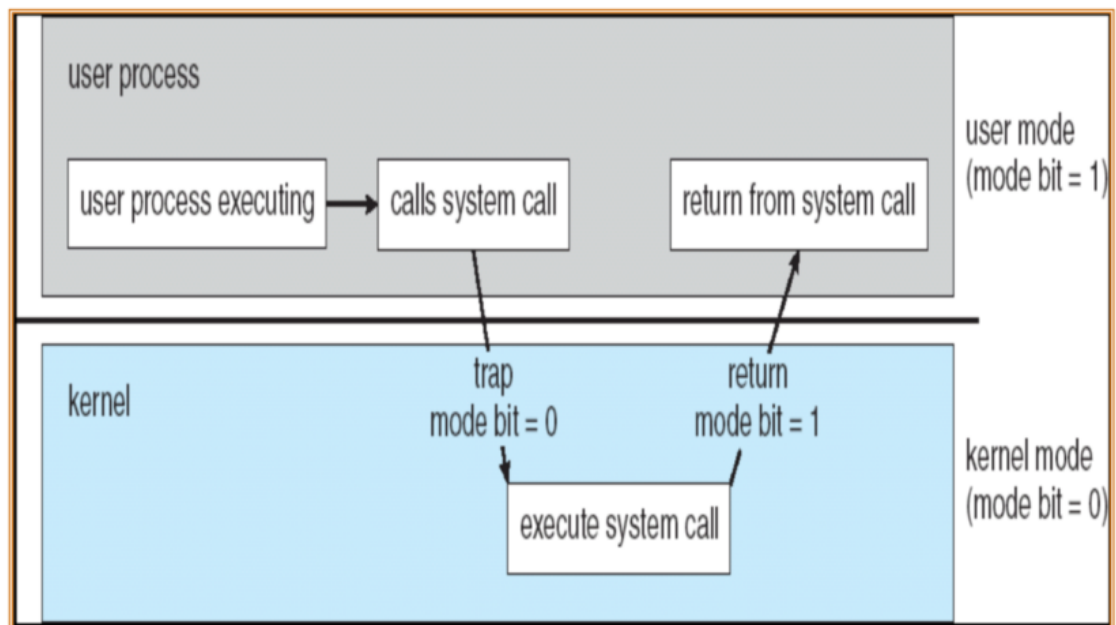
[Chapter 4 Threads](#)

[Chapter 5 CPU Scheduling](#)

[Chapter 6 Process Synchronization](#)

Chapter 1 Introduction

- System view
 - Resource allocator
 - Control program
- Dual-Mode Operation
 - User mode
 - Kernel mode
 - privileged instruction
 -



- Hardware
- **CPU protection**
 - **timer**
 - **time sharing**
- **memory protection**
 - **Base register**
 - **Limit register**
- **I/O protection**

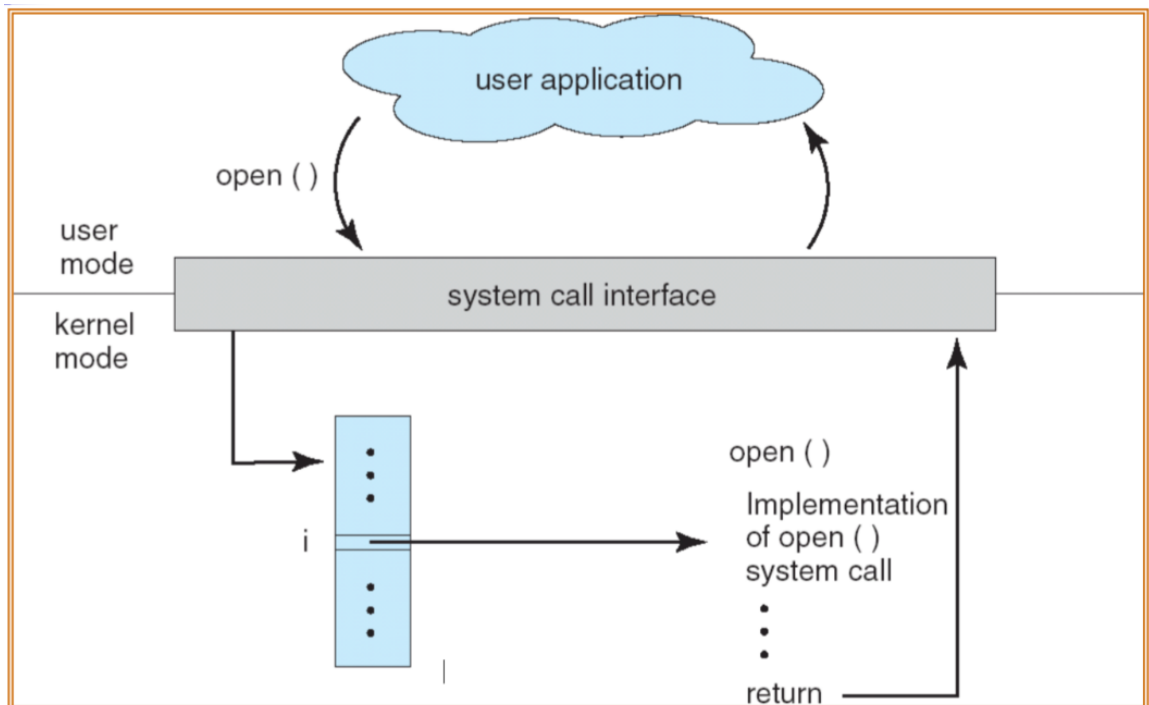
- **all I/O instruction are privilege instructions**
- Development of OS
 - mainframe systems
 - NO OS
 - **batch systems**
 - **multiprogramming systems**
 - **time sharing systems**
 - desktop systems
 - multiprocessor systems
 - distributed systems
 - clustered systems
 - real-time systems
 - handheld systems
 - 现代操作系统的特征
 - **并发性Concurrence**
 - **共享性Sharing**
 - **虚拟性Virtual**
 - **异步性Asynchronism**
 - 提高CPU利用率，充分发挥并发性：**程序之间、设备之间、设备与CPU之间均并发**
 - Pr:

批处理系统、多道程序系统和分时系统的技术特性

Chapter 2 Operating-System Structures

- 功能和服务的差别：
 - 对内：自行实现
 - 对外：可以调用其他功能代为实现
- common function of OS
 - process management
 - process synchronization
 - process communication
 - deadlock handling
 - (分布式)
 - main memory management
 - secondary-storage management
 - file management
 - I/O system management
- Operating System Services(Services for **helping users**)
 - Program execution
 - I/O operations
 - File-system manipulation
 - Communications

- Error detection
- Resource allocation
- Accounting(审计)
- Protection
- Operating System Interface
 - Interface to programs
 - **System calls**
 - System-call interface(SCI)
 - Application Programming Interface(API)
 - managed by runtime support library

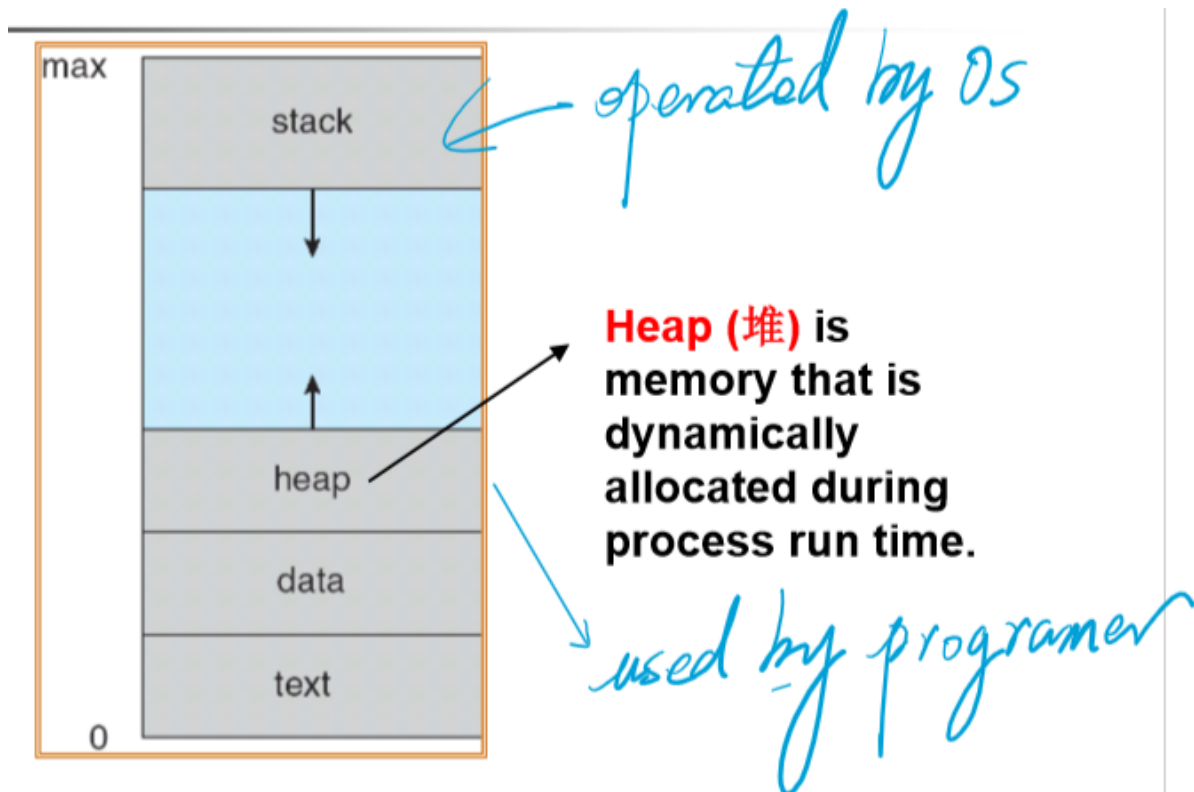


- Types of System calls
 - Process control
 - File management
 - Device management
 - Information maintenance
 - Communications
- **PR. Why do user use APIs rather than system calls directly?**
- **ANS.**
 1. 跨平台能力 (提供相同的API封装) 移植性好
 2. 模块化封装, 可维护性好
 3. 简化了程序编写
 4. 提高了执行效率
- Operating System Structure
 - Simple structure
 - Layered structure

- virtual machines
- Microkernel structure
 - Benefis
 - easier to extend
 - easier to port
 - more reliable
 - more secure
- Modules
- PR:设计操作系统时采用的模块化内核方法和分层方法在那些方面类似？ 那些方面不同？
- Operating system design and implementation
- 小结
 - 操作系统概念（管理资源、支持程序运行、方便用户使用的**程序集**）
 - 操作系统的基本目标：**方便性和高效性**
 - 引导程序：**中断、中断处理程序、中断向量**
 - 储存结构：内存（**小、易失**）二级储存（**大、非易失**）、分层结构
 - I/O结构：设备控制器（本地缓冲）、DMA
 - 硬件保护：**双重模式操作、特权指令、I/O保护、内存保护、CPU保护**
 - 操作系统的发展：e.g: 多道程序设计
 - 操作系统的功能：进程（CPU）管理、内存管理、磁盘管理、文件管理、I/O管理、**用户接口**
 - 操作系统服务：**程序执行、I/O操作、文件系统操作、通信、错误检测与处理**、资源分配、统计、保护
 - 操作系统接口：用户接口（CLI、GUI）、程序接口（**系统调用（参数传递、类型）**）、SCI、API
 - 操作系统结构

Chapter 3 Process

- Process
 - test section(program code)
 - **program counter**
 - **contents of the processor's registers**
 - Heap-stack
 - data section
 -
 -



◦ Characteristic of process

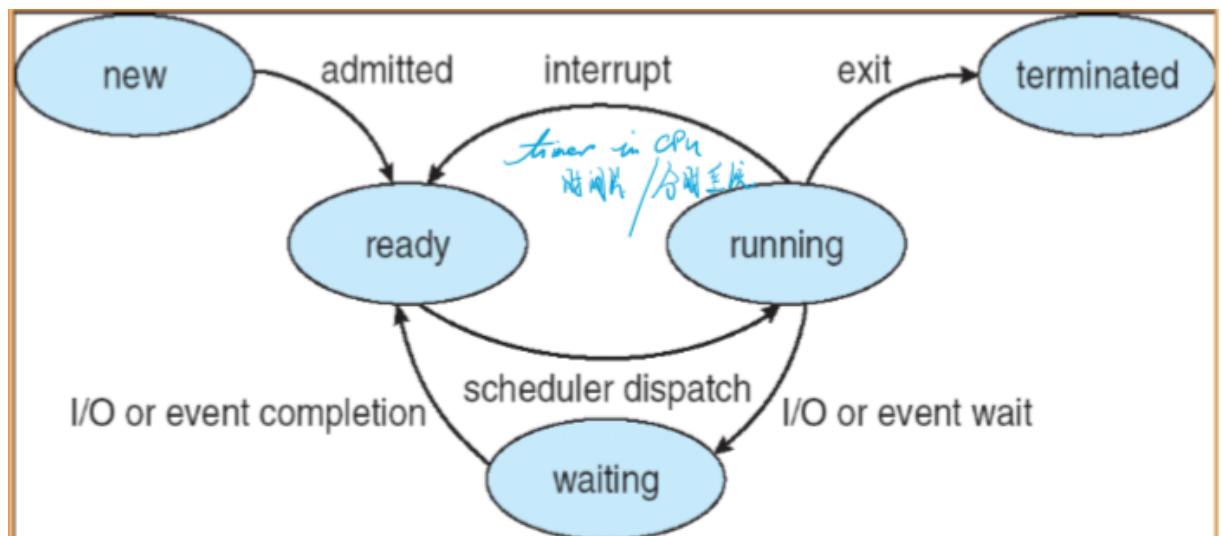
- Dynamic动态性
- Independency独立性
- Concurrency并发性
- Structure结构化

◦ PR.进程和程序是两个密切相关的概念，请阐述他们之间的区别和联系

◦ Process state

◦

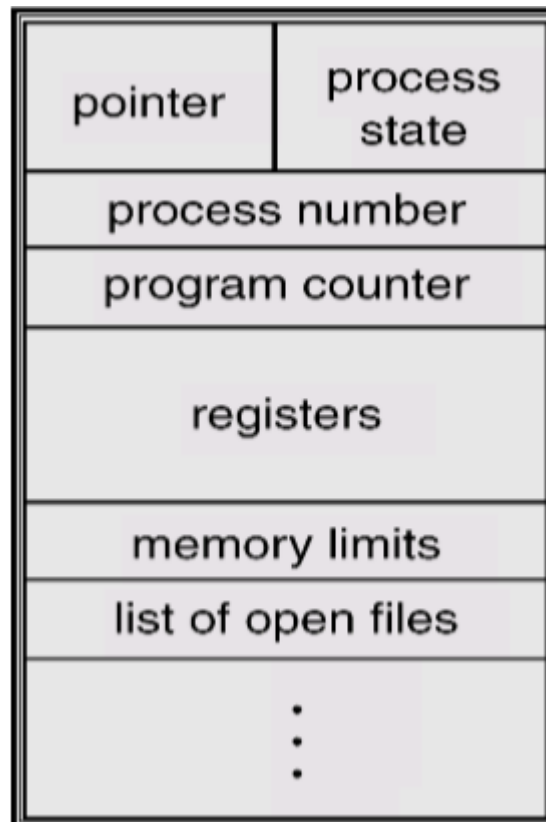
◦



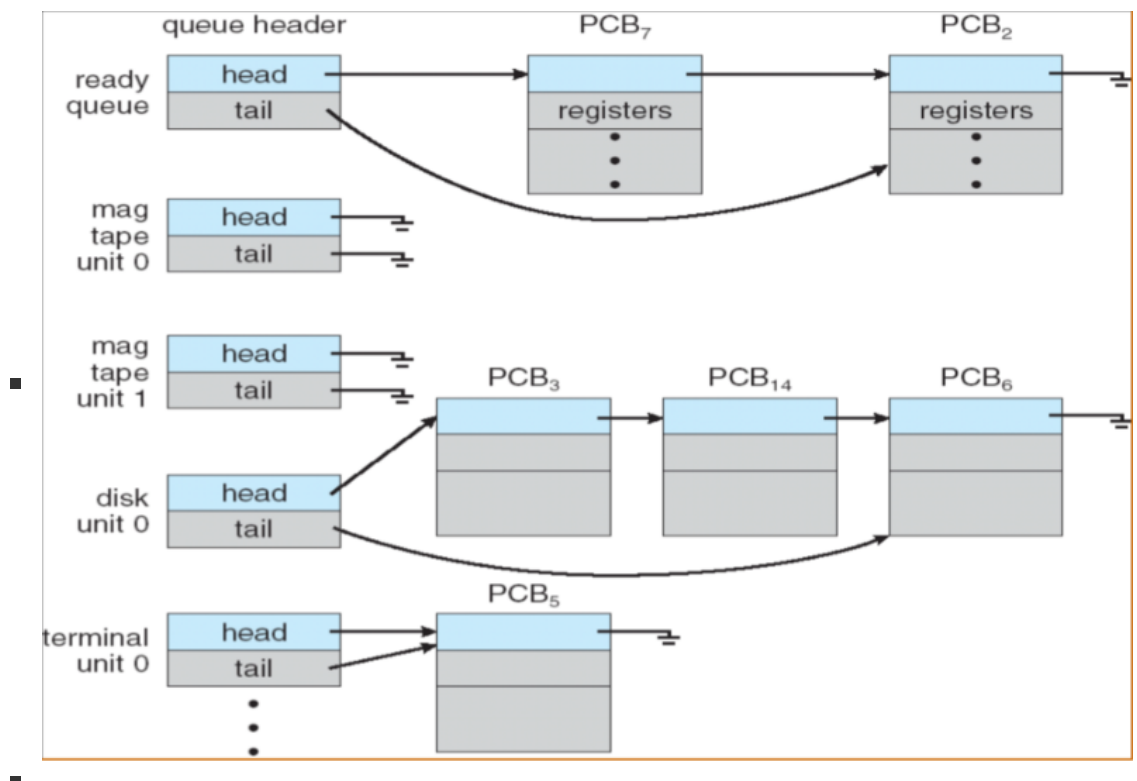
◦ Process control block(PCB)

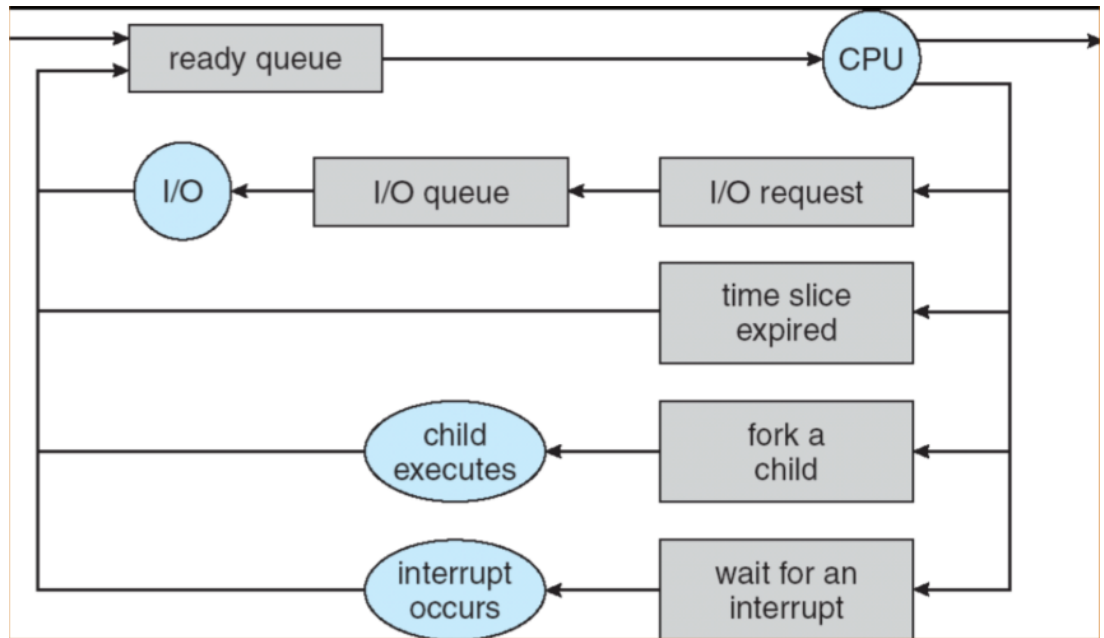
◦

◦

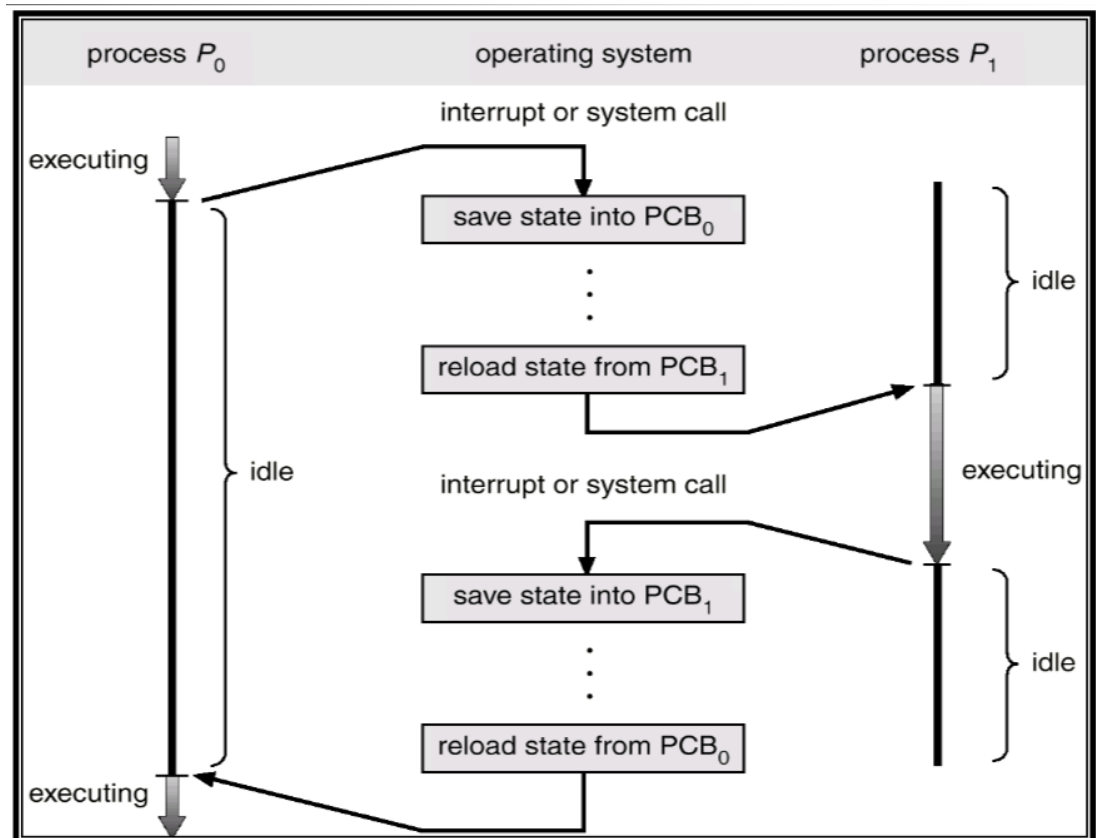


- Process scheduling queues
 - Job queue (in main memory)
 - Ready queue
 - device queues
 - process migration between the various queues






- Schedulers
 - Long-term scheduler(秒级、分钟级，作业调度)
 - Short-term scheduler(毫秒级，CPU调度)
 - Medium-term scheduler(swapping)
- I/O bound process
- CPU bound process
- Context switch
 - The **context** of a process is represented in **PCB** of the process and includes the values of CPU registers.
 - 保存执行后的上下文信息
 - 上下文切换会带来开销
 - 尽量减少上下文切换以减少开销
 -



- Operation on Process

- Process creation

- child process(unique process identifier(int)), tree of process
 - resource sharing
 - parent and children shall all resources
 - children share subset of parent's resources
 - parent and child share no resources
 - Execution
 - Parent and children execute concurrently
 - Parent waits until children terminate
 - Address space
 - child duplicate of parent
 - child has a program loaded into it (new text section)
 - UNIX examples
 - `fork()` :create new process
 - `exec()` :used after a fork to replace the process's memory space with a new program
 -



```

1  pid = fork();
2  if(pid<0) /* error occured */
3  {
4      printf(stderr,"Fork failed");
5      exit(-1);
6  }
7  else if(pid==0) /* child process */
8  {
9      execlp("/bin/ls","ls",NULL);
10 }
11 else /* parent process */
12 {
13     wait(NULL); /* wait for child process to finish */
14     printf("Child complete");
15     exit(0);
16 }
  
```

■ Process Termination

- `exit()` process executes last statement and asks the operating system to delete it
 - output data from child to parent (via `wait`)
 - Process's resources are deallocated by OS
- `abort()` parent may terminate execution of children process
 - child has exceeded allocated resources
 - task assigned to child is no longer required
 - parent is exiting *not all of the operation system supports **Cascading termination**(级联终止)

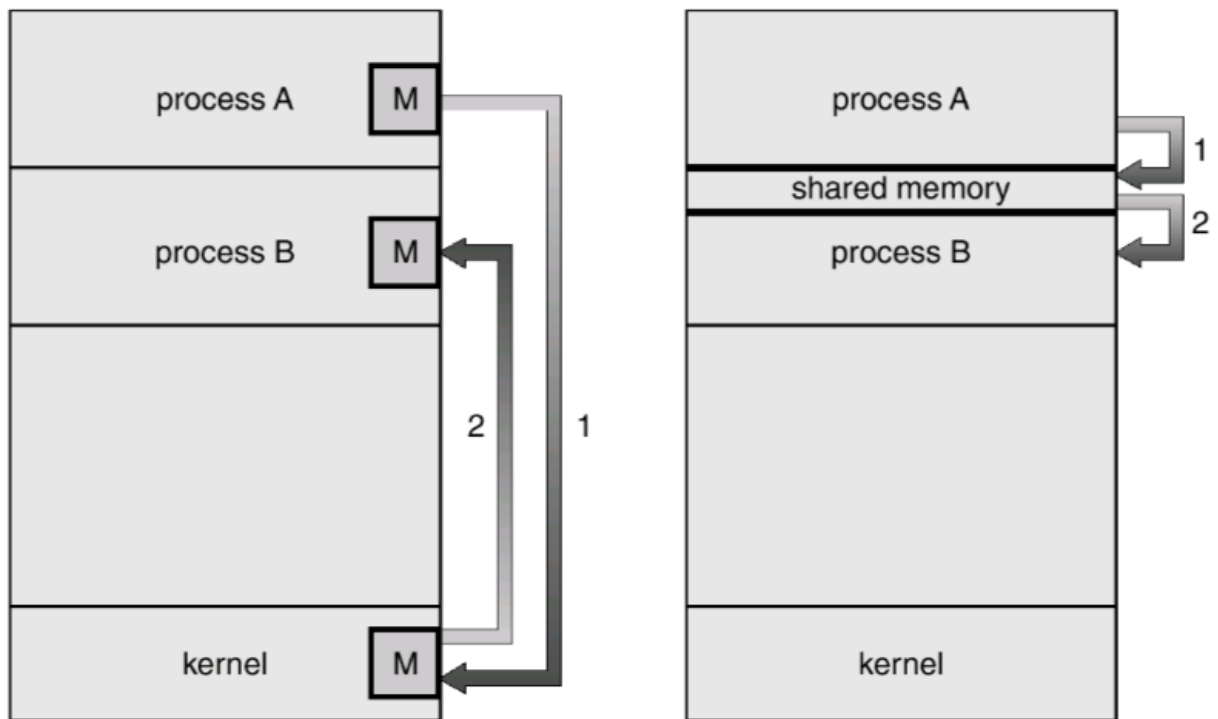
• InterProcess Communication(IPC)

- **Independent** process cannot affect or be effected by the execution of another process
- **Cooperating** process can affect or be effected by the execution of another process

■ Advantages

1. Information sharing
2. Computation speed-up
3. Modularity
4. Convenience

- **Shared memory & Message passing**

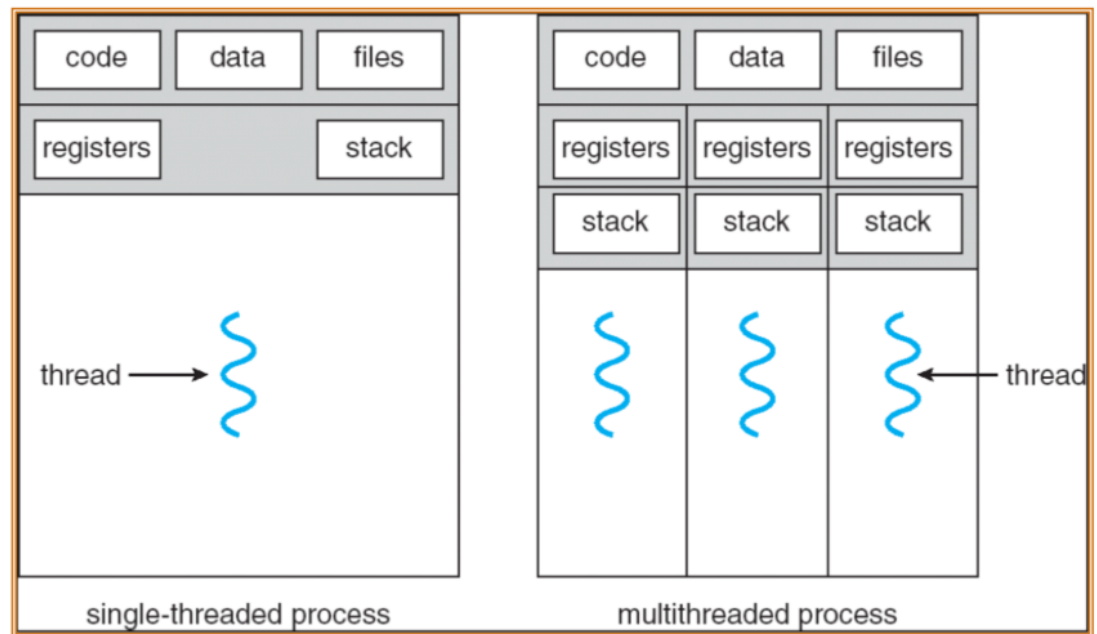


- Shared-memory Systems
 - requiring communication process to establish a region of shared memory
 - a shared memory region resides in the address space of the process creating the shared memory segment
 - the processes are responsible for ensuring that they are not writing to the **same location simultaneously**
 - Producer-Consumer Problem
- Message-passing Systems
 - MPS has two operations
 - send()
 - receive()
 - communication link
 1. link may be unidirectional or bidirectional
 2. a link may be associated with many processes
 - direct communication
 - send(P,message) send a message to process P
 - receive(Q,message) receive a message from process Q
 - indirect communication
 - mailboxes
 - each mailbox has a unique id
 - two processes can communicate only if they share a mailbox
 - Operations
 1. create a new mailbox
 2. send and receive messages through mailbox
 3. destroy a mailbox


- Synchronization
 - Blocking: synchronous
 - Non-blocking: asynchronous
- Buffering
 - **Zero capacity** sender must wait for receiver
 - **Bounded capacity** finite length of n messages, sender must wait if link full
 - **Unbounded capacity** infinite length, sender never blocks
- Communication in Client-Server System
 - **Sockets**
 - **Remote Procedure Calls**
 - **Remote Method Invocation (Java)**

Chapter 4 Threads

- Multithreading Models
 - A thread is a flow of control within a process
 - thread is a **basic** unit of CPU execution (known as LightWeight Process(LWP))
 - process (HeavyWeight process(HWP)) has a **single** thread of control
 - multithreaded process contains several **different** flows of control within the **same** address space
 - Thread
 - has
 - thread ID
 - program counter
 - register set
 - stack
 - share
 - code section
 - data section
 - other OS resources(file and signals)




- Benefits
 - responsiveness
 - resource sharing
 - economy(low cost in overhead of creating and context-switch)
 - Utilization of multiprocessor architectures
- User Threads
 - user threads are supported above the kernel. The kernel is **not** aware of user threads
 - Library provides all support for thread creation, termination, joining and scheduling
 - more efficient(no kernel intervention)
 - if one thread is blocked, every other threads of the same process are also blocked(containing process is blocked)



- Kernel Threads

- kernel threads are usually **slower** than the user threads
- blocking one thread will **not** cause other threads of the same process to block
- the kernel can schedule threads on different processors(in a multiprocessor environment)

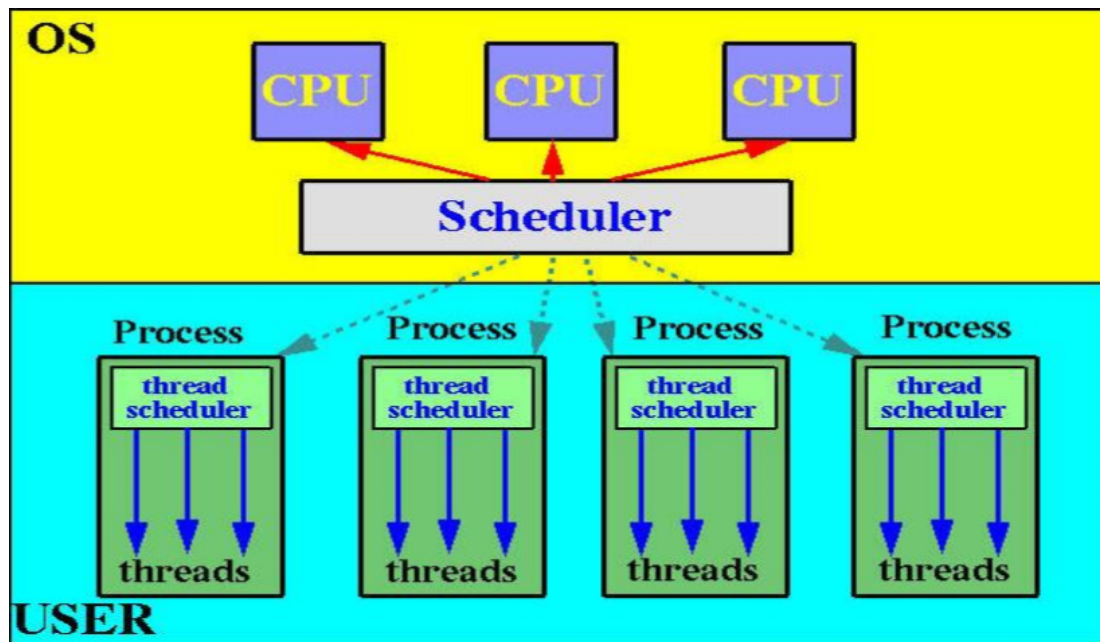


Pr.

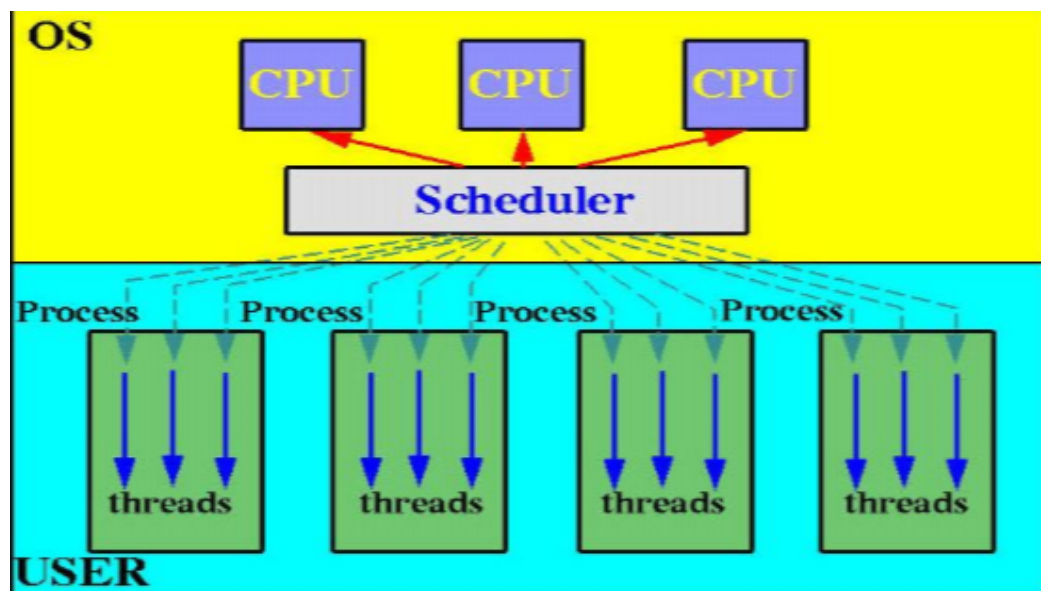
1. 进程和线程之间的区别和联系
2. 用户级线程和内核级线程的区别

■ Multithreading models

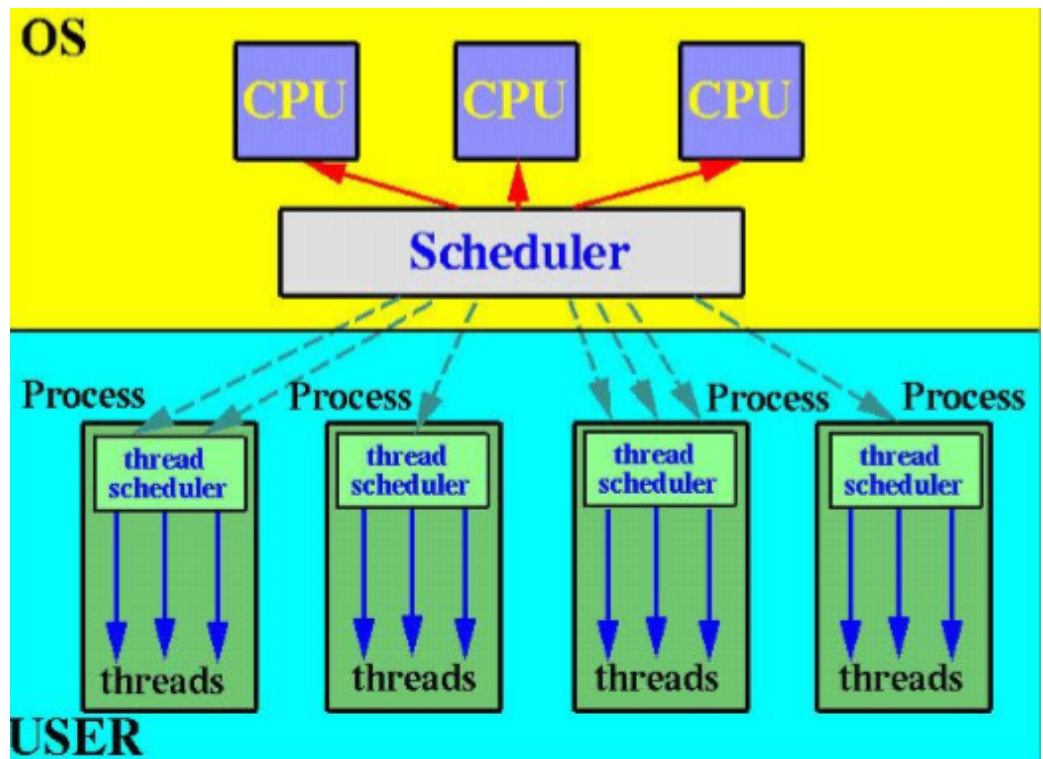
- many to one
 - only one thread in the one process can access the kernel at a time
 - true concurrency is not gained



- one to one
 - each user-level thread maps to kernel thread
 - providing more concurrency
 - restricting the number of threads supported by the system



- many to many
 - allow many user level threads to be mapped to many kernel threads



- Thread Libraries

- status
-

```

1  int pthread_create(tid,attr,function,arg);
2  /*
3   * pthread_t *tid
4   *   handle of created thread
5   * const pthread_attr_t *attr
6   *   attributes of thread to be created
7   * void *(*function)(void)
8   *   function to be mapped to thread
9   * void * arg
10  *   single argument to function
11  */
12  int pthread_join(tid,val_ptr);
13  /*
14  * pthread_t *tid
15  *   handle of joinable thread
16  * void ** var_ptr
17  *   exit value return by joined thread
18  */
19  void pthread_exit(void *status);
20  int pthread_cancel(pthread_t thread); //terminated immediately
21  int pthread_kill(pthread_t thread,int sig);

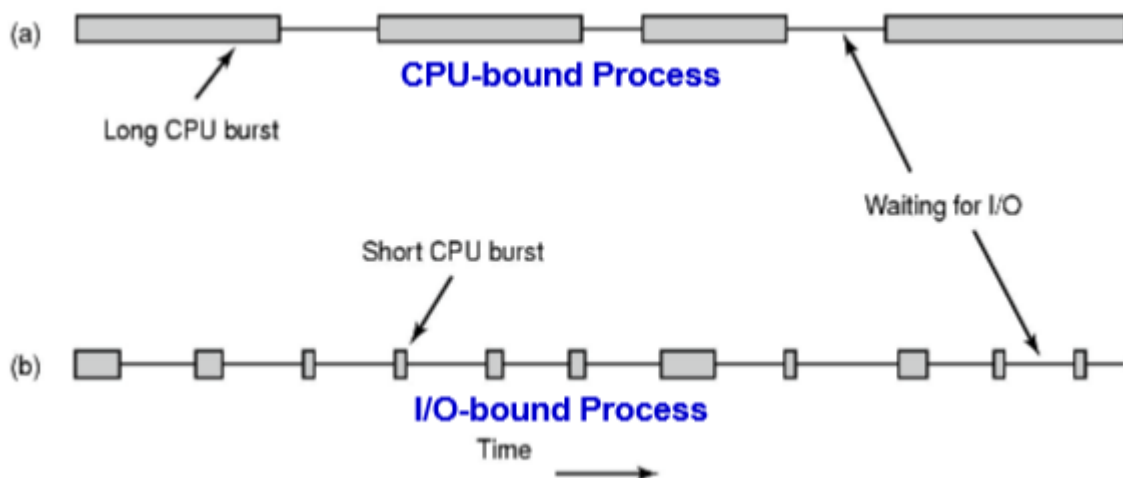
```

- CreateThread
- GetCurrentThreadId
- GetCurrentThread
- SuspendThread/ResumeThread

- ExitThread
- TerminateThread
- GetExitCodeThread
- GetThreadTimes
- Threading Issues
- Operating System Examples
- //TODO 关于线程的实现
- **Pr.**
 - 信号机制和中断机制的异同
- Thread Pools
 - advantages
 - faster to service a request(save the time to create new thread)
 - allow the number of threads in the application to be bound to the size of the pool
- Thread specific data
 - threads belonging to a process share the data of the process
 - allows each thread to have its own copy of data
 - when using a thread pool, each thread may be assigned a unique identifier
- Scheduler activations
- **upcalls**

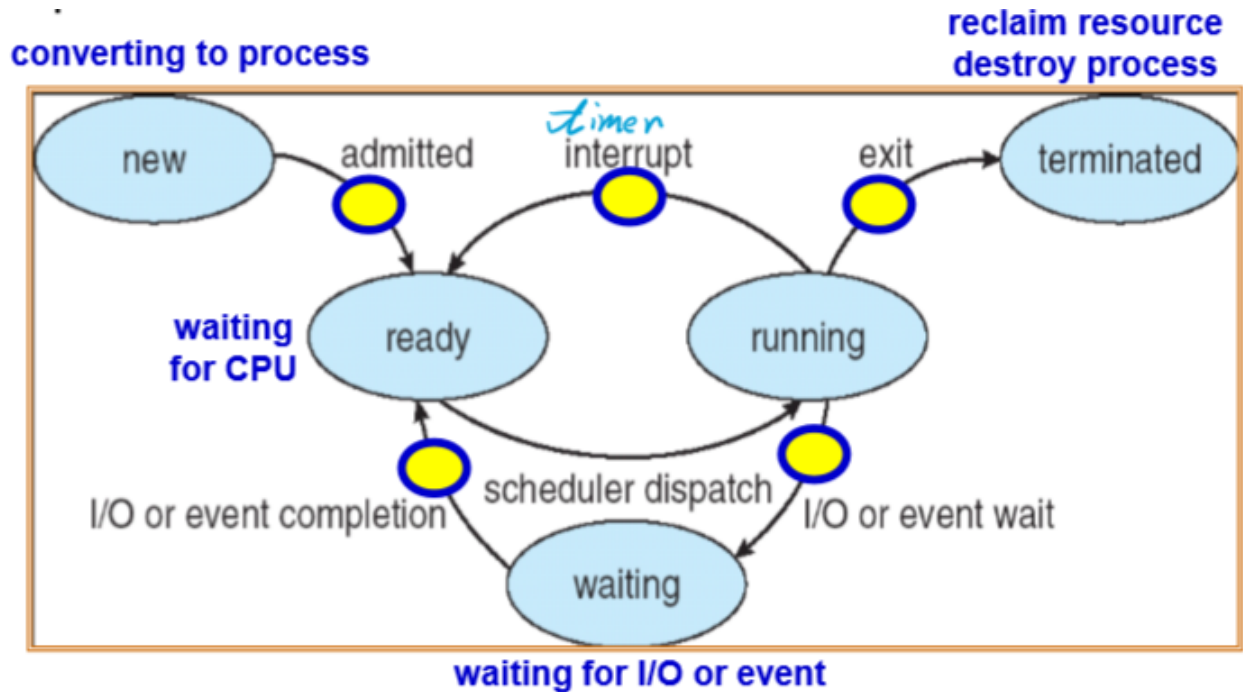
Chapter 5 CPU Scheduling

- Maximum CPU utilization obtained with multiprogramming
- The success of CPU scheduling depends on an property of processes: **CPU-I/O Burst Cycle**
 - process execution consists of a **cycle** of CPU execution and I/O wait.
- CPU-bound
 - a few very long CPU bursts
- I/O-bound
 - many short CPU bursts



- When the CPU is idle, the OS must select another ready process to run

- This selection process is carried out by the **short-term scheduler**
- The CPU scheduler selects a process from **the ready queue** and allocates the CPU to it
- There are many ways to organize the ready queue(e.g. FIFO)



- Circumstances that scheduling may take place
 - A process switches from the running state to the terminated state(finished)
 - A process switches from the running state to the wait state(e.g. IO operation)
 - ↑主动操作↑ 非抢占式调度
-
- ↓被动中止↓ 抢占式调度 → 同步机制
 - A process switched from the running state to the ready state(e.g. a interrupt occurs)
 - A process switches from the wait state to the ready state(e.g. I/O completion)
 - A process switches from the new state to ready state(e.g. a higher priority process ready)
 - Preemptive(抢占式)
 - cost associated with access to **shared data**
 - When the kernel is in its **critical** section modifying some important data .
 - special attention to situation
 - Non-preemptive
 - scheduling occurs when a process **voluntarily terminates**(主动结束) (case1)or enters the wait state(case2)
 - simple but very inefficient

Pr.

对于计算中心, 抢占式调度和非抢占式调度哪一种比较适合

- Dispatcher(调度) module
 - switching context

- switching to user mode
 - jumping to the proper location in the user program to restart that program
- Dispatch latency
 - the dispatcher should be as fast as possible
- Scheduling criteria
 - CPU utilization
 - keep the CPU as busy as possible
 - lightly | 40% | - | 90% | heavily
 - Throughput(吞吐)
 - higher throughput means more jobs get done

吞吐量和CPU利用率有相关性但并没有直接关系

- Turnaround time
 - The time period from job submission to completion is the turnaround time

$$\begin{aligned}
 t_{\text{turnaround}} = & \\
 & t_{\text{waitingTimeBeforeEnteringTheSystem}} + \\
 & t_{\text{waitingTimeInTheReadyQueue}} + \\
 & t_{\text{waitingTimeInAllOtherEvents}} + \\
 & t_{\text{timeTheProcessActuallyRunningOnTheCPU}}
 \end{aligned}$$

- Waiting time
 - time in ready queue
- Response time
 - the time from the submission of a request
- Optimization Criteria
 - MAX CPU utilization
 - MAX throughput
 - MIN turnaround time(average)
 - MIN waiting time
 - MIN response time

- 为什么需要CPU调度

大多数任务是CPU和I/O交替使用,

导致CPU和I/O至少有一个空闲,

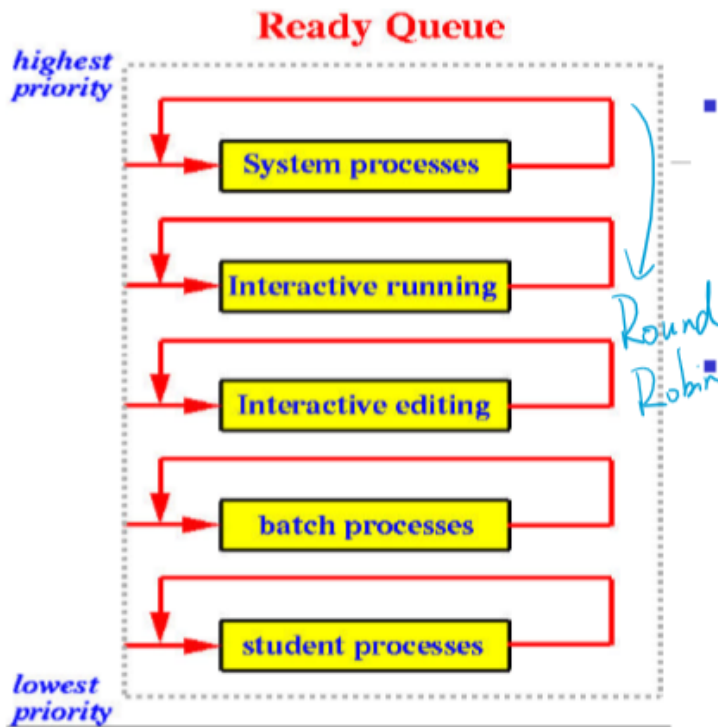
通过调度让需要执行I/O的任务去执行I/O。

把CPU给需要CPU的任务运行。

- **Scheduling Algorithms**

- First-Come-First-Served Scheduling (FCFS)
 - can easily implemented using a queue
 - not preemptive
 - convoy effect (护航效应)
 - troublesome for time-sharing systems
- Short-Job-First Scheduling (SJF)
 - sorted in next CPU burst length
 - can be nonpreemptive and preemptive


- **minimum average waiting time for a given set of process**
 - predict CPU burst: exponential averaging
 - long jobs may meet **starvation!!!**
- Priority Scheduling
 - each process has a **priority**
 - priority may be determined internally or externally
 - internal priority
 - time limits
 - memory requirement
 - number of files
 - etc.
 - external priority
 - importance of the process (not controlled by the OS)
 - starvation/Indefinite block
 - a lower priority may never have a chance to run
 - Aging
 - gradually increase the priority of process what wait in the system for a long time
- Round_Robin Scheduling (RR)(轮询)
 - designed for time-sharing systems
 - each process is assigned a time quantum/slice
 - If the process uses CPU for less than one time quantum, it will release the CPU voluntarily (主动退出)
 - when one time quantum is up , that process is preempted by the scheduler and moved to the tail of the list
 - Typically, higher average time than SJF, better response time
 - time quantum is too large → FCFS
 - time quantum is to small → processor sharing (并发)
 - *shorter time quantum means more context switches*
 - in general, 80% of the CPU bursts should be shorter than the time quantum
 -
- Multilevel Queue Scheduling (多级队列)
 - partitioned into separate queues
 - foreground (interactive)
 - background (batch)
 - Each process is assigned permanently to one queue based on some properties of the process
 - Each queue has its own scheduling algorithm
 - foreground - RR
 - background -FCFS



- A process P can run only if all queues above the queue that contains P are empty.

■ When a process is running and a process in a higher priority queue comes in, the running process is preempted.

- Scheduling must be done between the queues
 - Fixed priority scheduling (possibility of starvation)
 - Time slice
 - each queue gets a certain amount of CPU time which it can schedule amongst its processes
- Multilevel Feedback Queue Scheduling
 - allows process to move between queues
 - aging can be implemented this way
 - If a process use more/less CPU time, it is moved to a queue of lower/higher priority → I/O/CPU-bound process will be in higher/lower priority queues
 - exp



- number of queues
- scheduling algorithms for each queue
- method used to determine when to upgrade a process
- method used to determine when to demote a process
- method used to determine which queue a process will enter when that process needs service
- Multiple-Processor Scheduling
 - Homogeneous(同构) processors
 - Load balancing
 - push migration
 - pull migration
 - Asymmetric multiprocessing (非平衡处理)
 - only on processor accesses the system data
 - alleviating(降低) the need for data sharing
 - Symmetric multiprocessing (SMP)
 - two processors do **not** choose the same process
 - Processor Affinity (侵核)
 - most SMP systems **try** to avoid migration of processes from one processor to another
 - Soft/Hard Affinity (执行过程中可以/不可以侵核)
- Real-Time Scheduling
 - Hard real-time systems
 - the scheduler either **admits** a process and guarantees that the process will complete on-time, or **reject** the request (resource reservation)
 - secondary storage and virtual memory will cause unavoidable delay
 - Hard real-time systems usually have special software on special hardware
- Soft real-time systems
 - easily doable(可行) within a general system
 - may cause unfair resource allocation and longer delay(starvation) for noncritical tasks.
 - the CPU scheduler must **prevent aging** to occur(critical tasks may have lower priority)

- **The dispatch latency must be small**
- Priority Inversion
 - a high-priority process needs to access the data that is currently being accessed by a low-priority process → The high-priority process is blocked by the low-priority process
 - priority-inheritance protocol
- Thread Scheduling
 - User-level threads
 - thread library
 - Kernel-level threads
 - scheduled by OS
 - user-level threads must ultimately be mapped to an associated kernel-level thread
 - Local scheduling → User-level Thread
 - Process-contention Scope (PCS)
 - Global Scheduling → Kernel-level Thread
 - System-contention Scope (SCS)
- Algorithm Evaluation
 - Deterministic modeling (Analytic evaluation) 确定情况下 的情形证明
 - Queueing models 队列模型
 - Simulations 仿真
 - Implementation 证明
 - 从上往下证明力越强，越难证明
- Operating System
 - Scheduling threads using **preemptive** and **priority-based** scheduling algorithms (Real time, system, time sharing, interactive)
 - The default scheduling class for a process is time sharing (multilevel feedback queue)

Chapter 6 Process Synchronization

- Bounded-buffer

```

1  //Shared data
2  #define BUFFER_SIZE 10
3  typedef struct
4  {
5      //...
6  } item;
7  item buffer[BUFFER_SIZE];
8  int in = 0;
9  int out = 0;
10 int counter = 0;
11
12 //Producer process
13 item nextProduced;
14 while(1)

```

```

15 {
16     while(counter == BUFFER_SIZE);
17     //do nothing
18     buffer[in] = nextProduced;
19     in = (in + 1) % BUFFER_SIZE;
20     counter++;
21 }
22
23 //Consumer process
24 item nextConsumed;
25 while(1)
26 {
27     while(counter == 0)
28         //do nothing
29     nextConsumed = buffer[out];
30     out = (out + 1) % BUFFER_SIZE;
31     counter--;
32 }

```

- **Atomic operation**


- counter++
- counter—

- **Race condition**

- two or more processes/thread access and manipulate the same data concurrently
- the outcome of the execution depends on the particular order in which the access takes place
- To prevent race conditions, concurrent processes must be synchronized

- **The Critical-Section Problem**

- Each process has a code segment, called critical section
- **Problem:** ensure that when one process is executing in its critical section, no other process is allowed to execute in its critical section
- The critical-section problem is to design a protocol that processes can use to cooperate



critical section must run in a mutually exclusive way.

- **Solution to Critical-Section Problem**

- Mutual Exclusion (互斥、忙等) → 防止冲突
- Progress (空闲让进) → 进展性

- Bounded Waiting (有限等待) → 进展性
 - 防止饥饿, 让权等待, 多CPU: 死锁
- **the solution cannot depend on relative speed of processes and scheduling policy**
- Mutual Exclusion
- Bakery Algorithm

```

1  //shared data
2  boolean choosing[n];    //false
3  int number[n];          //0
4  do
5  {
6      choosing[i] = true;
7      number[i] = max(number[0], number[1], ..., number[n-1]) + 1;
8      choosing[i] = false;
9      for(j = 0; j < n; ++j)
10     {
11         while(choosing[j]);
12         while((number[j] != 0) && ((number[j], j) < (number[i], i)));
13     }
14     //critical section
15     number[i] = 0;
16     //remainder section
17 }while(1)

```

- Interrupt Disabling
 - disable interrupts → critical section → enable interrupts
 - When interrupts are disabled, no context switch will occur in a critical section
 - Infeasible in a multiprocessor system because all CPUs must be informed
 - Some feature that depend on interrupts (e.g. clock) may not work properly
- Mutual Exclusion (互斥锁)
 - TestAndSet

```

1  boolean TestAndSet(boolean &target)
2  {
3      boolean rv = &target;
4      &target = true;
5      return rv;
6  }

```



```

1 //shared data
2 boolean lock = false;
3 //Process P
4 do
5 {
6     while(TestAndSet(lock));
7     //critical section
8     lock = false;
9     //remainder section
10 }

```

- Swap
 - **atomically** swap two variables

```

1 void Swap(boolean &a,boolean &b)
2 {
3     boolean temp = &a;
4     &a = &b;
5     &b = temp;
6 }

```

```

1 //Global shared data
2 boolean lock; //false
3 //Local variable for each process
4 boolean key;
5 Process Pi
6 do
7 {
8     key = true;
9     while(key == true)
10 {
11     swap(lock,key);
12 }
13 //critical section
14 lock = false;
15 //remainder section
16 }

```

- Semaphores

```

1 wait(S)
2 {
3     while(S <= 0);
4     --S;
5 }
6
7 signal(S)
8 {
9     ++S;
10 }

```

- Count semaphore
- Binary semaphore (mutex locks)
- busy waiting (Spinlock)
- block itself (阻塞方法, 使用PCB唤醒)
 - Define a semaphore as a record

```

1  typedef struct
2  {
3      int value;
4      struct process *L;  //waiting queue
5  }semaphore;

```

- block()
- wakeup(P)

```

1  wait(S)
2  {
3      S.value--;
4      if (S.value < 0)
5      {
6          //add this process to S.L;
7          block();
8      }
9  }
10 signal(S)
11 {
12     S.value++;
13     if(S.value <= 0)
14     {
15         //remove a process P from S.L;
16         wakeup(P);
17     }
18 }

```

- if the semaphore is negative, its magnitude is the number of process waiting on that semaphore
- Busy waiting has not been **completely** eliminated
- furthermore, we have limited busy waiting to the critical sections of the wait() and signal() operations
- Deadlock and Starvation

临界资源、同步关系

- Bounded-Buffer Problem

```

1  //Shared data
2  Semaphore full = 0, empty = n, mutex = 1;
3  do //Producer
4  {

```

```

5      //produce an item in nextP
6      wait(empty);
7      wait(mutex);
8      //add nextP to buffer
9      signal(mutex);
10     signal(full);
11 }while(1);
12
13 do //Consumer
14 {
15     wait(full);
16     wait(mutex);
17     //remove an item from buffer to nextC
18     signal(mutex);
19     signal(empty);
20     //consume the item in nextC
21 }while(1);

```

- Readers and Writers Problem

- Reader first
- Writer first

```

1  //Shared data
2  int readcount;
3  semaphore wrt = 1,mutex = 1;
4  int readcount = 0;
5  do
6  {
7      wait(wrt);
8      //writing
9      signal(wrt);
10 }while(1);
11 do //Error: 写者饥饿问题
12 {
13     wait(mutex);
14     readcount++;
15     if(readcount == 1)
16         wait(wrt);
17     signal(mutex);
18     //reading
19     wait(mutex);
20     readcount--;
21     if(readcount == 0)
22         signal(wrt);
23     signal(mutex);
24 }

```

- Dining-Philosophers Problem
- 过独木桥问题

```

1 //Shared data
2 int countA = 0; //A方向上已在独木桥上的行人数目
3 int countB = 0; //B方向上已在独木桥上的新人数目
4 semaphore MA = 1; //countA的互斥锁
5 semaphore MB = 1; //countB的互斥锁
6 semaphore mutex = 1; //实现互斥使用

```

■ A方向过桥

```

1 do
2 {
3     wait(MA);
4     countA++;
5     if (count == 1)
6     {
7         wait(mutex);
8     }
9     signal(MA);
10    //过桥
11    wait(MA);
12    countA--;
13    if(countA == 0)
14    {
15        signal(mutex);
16    }
17    signal(MA);
18 }while(1);

```

• Monitors (管程)

- High-level synchronization construct that allows the safe sharing of an abstract data type among concurrent processes


```

1 monitor monitor-name
2 {
3     shared variable declarations
4     procedure body P1()
5     {
6         //...
7     }
8     procedure body P2()
9     {
10        //...
11    }
12    //...
13    { //initialization code }
14 }

```

- no more than one process can be executing within a monitor
- when a process calls a monitor procedure and the monitor has a process running, the caller will be blocked outside the monitor


- Mutual exclusion is guaranteed with in a monitor



- Condition variables

- x, y

- $x.\text{wait}()$ means that the process invoking this operation is suspended until another process invokes $x.\text{signal}()$;
- $x.\text{signal}()$ operation resumes exactly one suspended process. If no process is suspended, the $\text{signal}()$ operation has no effect



Semaphores	Condition Variables
Can be used anywhere, but not in a monitor	Can only be used in monitors
wait() does not always block its caller	wait() always blocks its caller
signal() either releases a process, or increase the semaphore counter	signal() either releases a process ,or the signal is lost as if it never occurs
If signal() release a process, the caller and the release both continue	If signal() release a process, either the caller or the released continues, but not both

- 管程是公用数据结构，进程是私有数据结构
- 管程集中管理共享变量上的同步操作，临界区分散在每个进程中
- 管程管理共享资源，进程占用系统资源和实现系统并发性
- 管程被欲使用的共享资源的进程调用，管程和调用它的进程不能并发工作，进程之间能并发工作
- 管程是语言或操作系统的成分，不必创建或撤销，进程有生命周期，有创建有消亡

🚧正在施工中.....🚧